

32-bit RISC-V Processor

Design Specification

Author

Marco

Development Environment

HDL Verilog HDL

EDA Tool Vivado 2025.1

ISA RV32I + RV32M

Target Xilinx ZCU104

Contents

| | |
|---|-----------|
| 1 Abstract & Introduction | 3 |
| 2 RV32I & RV32M | 4 |
| 2.1 The RV32I & RV32M Instruction Table supported by this processor | 5 |
| 3 RISC-V Core | 7 |
| 3.1 Processor Core | 8 |
| 3.1.1 Datapath Descriptions for All Instructions | 8 |
| 3.2 Program Counter (PC) | 14 |
| 3.3 Register File / General-Purpose Register (GPR) | 14 |
| 3.4 Immediate Generator (ImmGen) | 15 |
| 3.5 Processor Controller | 15 |
| 3.6 Arithmetic Logic Unit (ALU) | 16 |
| 3.7 ALU Controller | 16 |
| 3.8 Branch Processing Unit (BPU) | 16 |
| 3.9 Load Data Unit (LDU) | 17 |
| 3.10 Data Memory | 17 |
| 3.11 Five-Stage Pipelined System | 17 |
| 3.12 Forwarding, Hazard and Flush Detection Unit | 18 |
| 3.12.1 Forwarding Detection | 18 |
| 3.12.2 Hazard Detection | 19 |
| 3.12.3 Flush Detection | 19 |
| 3.13 Branch History Table (BHT) & Branch Tag Buffer (BTB) | 20 |
| 3.13.1 BHT (Branch History Table) | 20 |
| 3.13.2 BTB (Branch Target Buffer) | 21 |
| 3.14 Control State Register (CSR) | 22 |
| 3.14.1 Supported CSR Registers | 22 |
| 3.15 Verification | 23 |
| 4 AXI-4 Lite Bus, BRAM and Cache | 24 |
| 4.1 AXI-4 Lite Interface with BRAM | 25 |
| 4.1.1 Required signals on an AXI4-Lite interface | 25 |
| 4.1.2 AXI4-Lite interface to BRAM | 26 |
| 4.2 Two-way set-associative Cache with AXI-4 Lite to BRAM | 27 |
| 4.2.1 Instruction Cache (Only Read Cache) | 27 |
| 4.2.2 Data Cache | 28 |

| | |
|---|-----------|
| 5 RISC-V Processor | 29 |
| 5.1 Full System Integration | 30 |
| 5.2 System-Level Verification | 31 |
| 5.3 Synthesis Results | 32 |

Chapter 1

Abstract & Introduction

This project is based on [Computer Organization Course](#) and [UC Berkeley CS 61C](#). The goal was to design a 32-bit pipelined RISC-V CPU supporting RV32I and RV32M, including Forwarding, Hazard Detection, Flush Detection, Dynamic Branch Prediction, and two Caches with AXI4-Lite Bus to BRAM. The final design was implemented and verified on FPGA.

Our development process follows a modular approach, beginning with the core execution unit and progressively expanding into a complete system-on-chip (SoC) architecture. The project is structured into four sequential phases :

1. **Core Processor Design:** Implement a 5-stage pipelined RISC-V CPU supporting both the RV32I (Base Integer) and RV32M (Integer Multiplication and Division) instruction sets, including Forwarding, Hazard Detection, Flush Detection, and Dynamic Branch Prediction.
2. **Bus Infrastructure:** Design an AXI4-Lite bus interface compliant with the ARM AXI protocol, followed by comprehensive system-level testing by interfacing the bus with BRAM (Block RAM).
3. **Memory Hierarchy:** Develop independent Instruction and Data Caches, and integrate them with the AXI4-Lite Bus and BRAM to verify cache coherency and proper timing behaviour.
4. **Full System Integration:** Integrate both caches into the pipelined RISC-V CPU, forming a complete processor connected to BRAM via the AXI4-Lite interconnect, and deploy the final design onto an FPGA.

The complete source code, RTL design files, testbenches, and verification scripts for this project are publicly available on GitHub:

<https://github.com/akira2963753/Pipelined-RV32-SoC>

Chapter 2

RV32I & RV32M

2.1 The RV32I & RV32M Instruction Table supported by this processor

Table 2.1: R-Type Instructions

| R-Type | Description | Funct7 | Rs2 | Rs1 | Funct3 | Rd | Opcode |
|--------|--------------------------------------|---------|-----|-----|--------|----|---------|
| ADD | $Rd = Rs1 + Rs2$ | 0000000 | Rs2 | Rs1 | 000 | Rd | 0110011 |
| SUB | $Rd = Rs1 - Rs2$ | 0100000 | Rs2 | Rs1 | 000 | Rd | 0110011 |
| SLL | $Rd = Rs1 \ll Rs2$ | 0000000 | Rs2 | Rs1 | 001 | Rd | 0110011 |
| SLT | $Rd = (Rs1 < Rs2)$ | 0000000 | Rs2 | Rs1 | 010 | Rd | 0110011 |
| SLTU | $Rd = (Rs1 < Rs2)$ | 0000000 | Rs2 | Rs1 | 011 | Rd | 0110011 |
| XOR | $Rd = Rs1 \oplus Rs2$ | 0000000 | Rs2 | Rs1 | 100 | Rd | 0110011 |
| SRL | $Rd = Rs1 \gg Rs2$ | 0000000 | Rs2 | Rs1 | 101 | Rd | 0110011 |
| SRA | $Rd = Rs1 \ggg Rs2$ | 0100000 | Rs2 | Rs1 | 101 | Rd | 0110011 |
| OR | $Rd = Rs1 Rs2$ | 0000000 | Rs2 | Rs1 | 110 | Rd | 0110011 |
| AND | $Rd = Rs1 \& Rs2$ | 0000000 | Rs2 | Rs1 | 111 | Rd | 0110011 |
| MUL | $Rd = (Rs1 \times Rs2)[31:0]$ | 0000001 | Rs2 | Rs1 | 000 | Rd | 0110011 |
| MULH | $Rd = (Rs1 \times Rs2)[63:32]$ | 0000001 | Rs2 | Rs1 | 001 | Rd | 0110011 |
| MULHSU | $Rd = (Rs1 \times Rs2[U])[63:32]$ | 0000001 | Rs2 | Rs1 | 010 | Rd | 0110011 |
| MULHU | $Rd = (Rs1[U] \times Rs2[U])[63:32]$ | 0000001 | Rs2 | Rs1 | 011 | Rd | 0110011 |
| DIV | $Rd = Rs1 / Rs2$ | 0000001 | Rs2 | Rs1 | 100 | Rd | 0110011 |
| DIVU | $Rd = Rs1[U] / Rs2[U]$ | 0000001 | Rs2 | Rs1 | 101 | Rd | 0110011 |
| REM | $Rd = Rs1 \% Rs2$ | 0000001 | Rs2 | Rs1 | 110 | Rd | 0110011 |
| REMU | $Rd = Rs1[U] \% Rs2[U]$ | 0000001 | Rs2 | Rs1 | 111 | Rd | 0110011 |

Table 2.2: I-Type Instructions

| I-Type | Description | Imm[11:0] | - | Rs1 | Funct3 | Rd | Opcode |
|--------|---------------------------------------|-----------|-------|-----|--------|----|---------|
| ADDI | $Rd = Rs1 + Imm$ | Imm[11:0] | - | Rs1 | 000 | Rd | 0010011 |
| SUBI | $Rd = Rs1 - Imm$ | Imm[11:0] | - | Rs1 | 001 | Rd | 0010011 |
| SLTI | $Rd = (Rs1 < Imm)$ | Imm[11:0] | - | Rs1 | 010 | Rd | 0010011 |
| SLTIU | $Rd = (Rs1 < Imm)$ | Imm[11:0] | - | Rs1 | 011 | Rd | 0010011 |
| XORI | $Rd = Rs1 \oplus Imm$ | Imm[11:0] | - | Rs1 | 100 | Rd | 0010011 |
| ORI | $Rd = Rs1 Imm$ | Imm[11:0] | - | Rs1 | 110 | Rd | 0010011 |
| ANDI | $Rd = Rs1 \& Imm$ | Imm[11:0] | - | Rs1 | 111 | Rd | 0010011 |
| JALR | $Rd = PC + 4, PC = Rs1 + Imm \& (-1)$ | Imm[11:0] | - | Rs1 | 000 | Rd | 1100111 |
| SLLI | $Rd = Rs1 \ll shamt$ | 0000000 | Shamt | Rs1 | 001 | Rd | 0010011 |
| SRLI | $Rd = Rs1 \gg shamt$ | 0000000 | Shamt | Rs1 | 101 | Rd | 0010011 |
| SRAI | $Rd = Rs1 \ggg shamt$ | 0100000 | Shamt | Rs1 | 101 | Rd | 0010011 |
| LB | $Rd[7:0] = M[Rs1+Imm]$ | Imm[11:0] | - | Rs1 | 000 | Rd | 0000011 |
| LH | $Rd[15:0] = M[Rs1+Imm]$ | Imm[11:0] | - | Rs1 | 001 | Rd | 0000011 |
| LW | $Rd = M[Rs1+Imm]$ | Imm[11:0] | - | Rs1 | 010 | Rd | 0000011 |
| LBU | $Rd[7:0] = M[Rs1+Imm]$ | Imm[11:0] | - | Rs1 | 011 | Rd | 0000011 |
| LHU | $Rd[15:0] = M[Rs1+Imm]$ | Imm[11:0] | - | Rs1 | 100 | Rd | 0000011 |
| CSRRW | $Rd = CSR[csr]; CSR[csr] = Rs1$ | CSR[11:0] | - | Rs1 | 001 | Rd | 1110011 |

Continued on next page

Table 2.2 – *Continued from previous page*

| I-Type | Description | Imm[11:0] | - | Rs1 | Funct3 | Rd | Opcode |
|--------|---------------------------------|-----------|---|-----|--------|----|---------|
| CSRRS | Rd = CSR[csr]; CSR[csr] = Rs1 | CSR[11:0] | - | Rs1 | 010 | Rd | 1110011 |
| CSRRC | Rd = CSR[csr]; CSR[csr] &= ~Rs1 | CSR[11:0] | - | Rs1 | 011 | Rd | 1110011 |

Table 2.3: S-Type Instructions

| S-Type | Description | Imm[11:5] | Rs2 | Rs1 | Funct3 | Imm[4:0] | Opcode |
|--------|--------------------------------|-----------|-----|-----|--------|----------|---------|
| SB | M[Rs1 + Imm][7:0] = Rs2[7:0] | Imm[11:5] | Rs2 | Rs1 | 000 | Imm[4:0] | 0100011 |
| SH | M[Rs1 + Imm][15:0] = Rs2[15:0] | Imm[11:5] | Rs2 | Rs1 | 001 | Imm[4:0] | 0100011 |
| SW | M[Rs1 + Imm] = Rs2 | Imm[11:5] | Rs2 | Rs1 | 010 | Imm[4:0] | 0100011 |

Table 2.4: B-Type Instructions

| B-Type | Description | Imm[12 10:5] | Rs2 | Rs1 | Funct3 | Imm[4:1 11] | Opcode |
|--------|-----------------------------------|--------------|-----|-----|--------|-------------|---------|
| BEQ | if(Rs1 == Rs2) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 000 | Imm[4:1 11] | 1100011 |
| BNE | if(Rs1 != Rs2) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 001 | Imm[4:1 11] | 1100011 |
| BLT | if(Rs1 < Rs2) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 100 | Imm[4:1 11] | 1100011 |
| BGE | if(Rs1 ≥ Rs2) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 101 | Imm[4:1 11] | 1100011 |
| BLTU | if(Rs1[U] < Rs2[U]) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 110 | Imm[4:1 11] | 1100011 |
| BGEU | if(Rs1[U] ≥ Rs2[U]) PC = PC + Imm | Imm[12 10:5] | Rs2 | Rs1 | 111 | Imm[4:1 11] | 1100011 |

Table 2.5: U-Type Instructions

| U-Type | Description | Imm[31:12] | Rd | Opcode |
|--------|-----------------------|------------|----|---------|
| LUI | Rd = Imm << 12 | Imm[31:12] | Rd | 0110111 |
| AUIPC | Rd = PC + (Imm << 12) | Imm[31:12] | Rd | 0010111 |

Table 2.6: J-Type Instructions

| J-Type | Description | Imm[20 10:1 11 19:12] | - | - | - | Rd | Opcode |
|--------|----------------------------|-----------------------|---|---|---|----|---------|
| JAL | Rd = PC + 4, PC = PC + Imm | Imm[20 10:1 11 19:12] | - | - | - | Rd | 1101111 |

Chapter 3

RISC-V Core

3.1 Processor Core

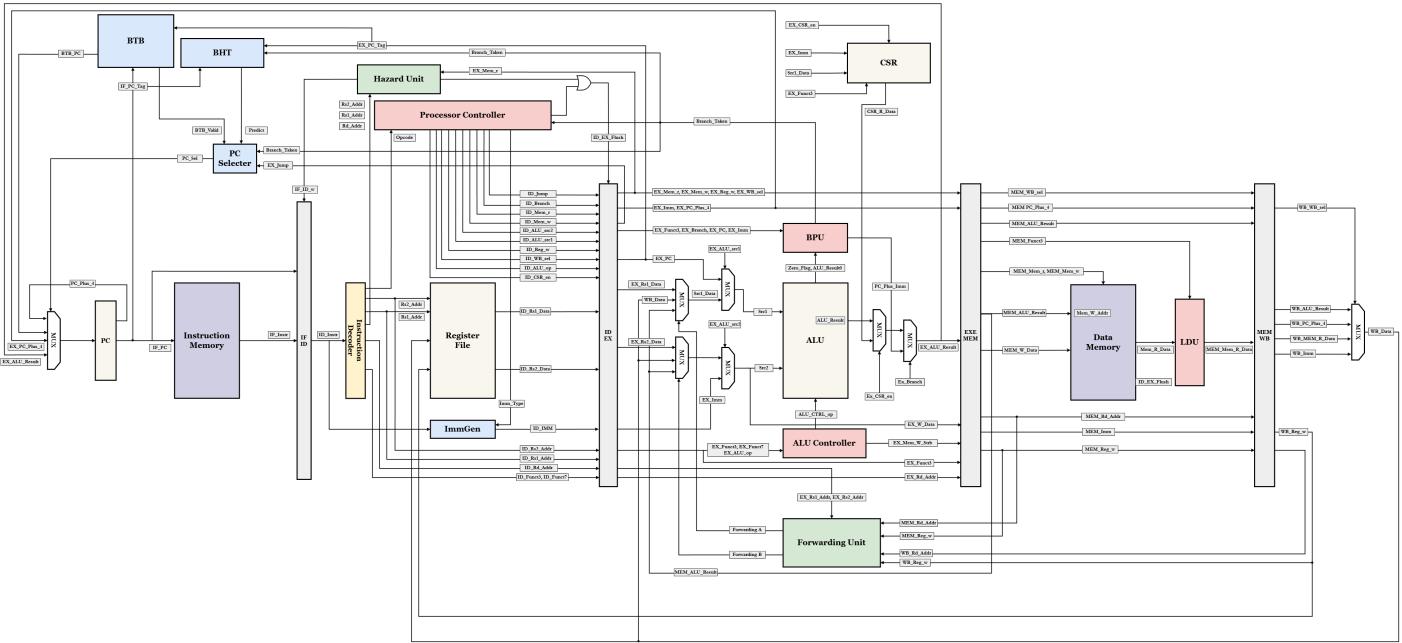


Figure 3.1: RISC-V Core System Framework Diagram Supporting RV32I and RV32M

3.1.1 Datapath Descriptions for All Instructions

Due to the large size of the datapath table, it will be presented across multiple pages.

Table 3.1: Instruction Operation Datapath Description

| Instruction | Operation Datapath Description |
|------------------|--|
| ADD Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully add the two data values. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SUB Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully subtract the two data values. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SLL Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \ll Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |

Continued on next page

Table 3.1 – *Continued from previous page*

| Instruction | Operation Datapath Description |
|---------------------|---|
| SLT Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to determine whether Rs1 is less than Rs2. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SLTU Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to determine whether Rs1 [U] is less than Rs2 [U]. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| XOR Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \oplus Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SRL Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \gg Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SRA Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \ggg Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| OR Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| AND Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \& Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| MUL Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \times Rs2$ and output the lower 32 bits. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| MULH Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \times Rs2$ and output the upper 32 bits. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| MULHSU Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \times Rs2 [U]$ and output the upper 32 bits. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |

Continued on next page

Table 3.1 – *Continued from previous page*

| Instruction | Operation Datapath Description |
|--------------------|--|
| MULHU Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 [U] \times Rs2 [U]$ and output the upper 32 bits. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| DIV Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 / Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| DIVU Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 [U] / Rs2 [U]$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| REM Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 \% Rs2$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| REMU Rd, Rs1, Rs2 | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 and EX_ALU_src2 = 0, allowing the ALU to successfully perform $Rs1 [U] \% Rs2 [U]$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| ADDI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform $Rs1 + Imm$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SUBI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform $Rs1 - Imm$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SLTI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to determine whether Rs1 is less than Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SLTIU Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to determine whether $Rs1 [U]$ is less than $Imm [U]$ (though Imm is still sign-extended first). Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| XORI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform $Rs1 \oplus Imm$. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |

Continued on next page

Table 3.1 – *Continued from previous page*

| Instruction | Operation Datapath Description |
|--------------------|--|
| ORI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| ANDI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 & Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| JALR Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm & (-1). Then, with EX_Jump = 1, PC_sel = 3, setting PC = EX_ALU_Result. In the WB stage, WB_WB_sel = 1, resulting in Rd = PC + 4. |
| SLLI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 << Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SRLI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 >> Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SRAI Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 >>> Imm. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| SB Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory write address. Then, W_STRB controls the byte write operation, writing Rs2[7:0]. |
| SH Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory write address. Then, W_STRB controls the half-word write operation, writing Rs2[15:0]. |
| SW Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory write address. Then, W_STRB controls the word write operation, writing Rs2[31:0]. |
| LB Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory read address. Then, the LDU controls and adjusts the output data, storing it into Rs2[31:0]. |

Continued on next page

Table 3.1 – *Continued from previous page*

| Instruction | Operation Datapath Description |
|--------------------|--|
| LH Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory read address. Then, the LDU controls and adjusts the output data, storing it into Rs2[31:0]. |
| LW Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory read address. Then, the LDU controls and adjusts the output data, storing it into Rs2[31:0]. |
| LBU Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory read address. Then, the LDU controls and adjusts the output data, storing it into Rs2[31:0]. |
| LHU Rd, Rs1, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 + Imm to calculate the memory read address. Then, the LDU controls and adjusts the output data, storing it into Rs2[31:0]. |
| BEQ Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |
| BNE Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |
| BLT Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |

Continued on next page

Table 3.1 – *Continued from previous page*

| Instruction | Operation Datapath Description |
|--------------------|--|
| BGE Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |
| BLTU Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |
| BGEU Rs1, Rs2, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm, which is also sent to the BPU to determine whether the branch is taken. If the branch is taken, it checks whether a prediction was made earlier. If the previous prediction was branch taken and the predicted address is correct, no jump is needed. Otherwise, PC_sel = 3, selecting EX_ALU_Result. |
| LUI Rd, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 0, EX_ALU_src2 = 1, allowing the ALU to successfully perform Rs1 = Imm << 12. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| AUIPC Rd, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm << 12. Then in the WB stage, with WB_WB_sel = 0, WB_ALU_Result is written back. |
| JAL Rd, Imm | After decoding in the ID stage, Forwarding A & B = 0, selecting EX_Rd_Data as the data for Rs1 and Rs2. EX_ALU_src1 = 1, EX_ALU_src2 = 1, allowing the ALU to successfully perform PC + Imm. Then, with PC_sel = 3, PC = PC + Imm. In the WB stage, WB_WB_sel = 1, resulting in Rd = PC + 4. |

3.2 Program Counter (PC)

The PC Module is responsible for managing the program counter updates and control, tracking the address of the instruction currently being executed by the CPU.

Table 3.2: PC Module I/O Ports

| I/O | Width | Port |
|--------|-----------|---------------|
| Input | 1 | CLK |
| Input | 1 | RST_N |
| Input | 2 | PC_sel |
| Input | 'PC_WIDTH | EX_ALU_Result |
| Input | 'PC_WIDTH | IF_PC_Plus_4 |
| Input | 'PC_WIDTH | BTB_PC |
| Input | 'PC_WIDTH | EX_PC_Plus_4 |
| Output | 'PC_WIDTH | IF_PC |

3.3 Register File / General-Purpose Register (GPR)

The Register File provides read and write functionality for 32 general-purpose registers (x0-x31), serving as data storage during CPU operations. The Register File features two read ports with direct read access, and one write port that writes on the rising edge of the clock. Additionally, the x0 register is hard-wired to zero.

Table 3.3: Register File I/O Ports

| I/O | Width | Port |
|--------|-------------|-------------|
| Input | 1 | CLK |
| Input | 1 | RST_N |
| Input | 1 | WB_Reg_w |
| Input | 'ADDR_WIDTH | ID_Rs1_Addr |
| Input | 'ADDR_WIDTH | ID_Rs2_Addr |
| Input | 'ADDR_WIDTH | WB_Rd_Addr |
| Input | 'DATA_WIDTH | WB_Rd_Data |
| Output | 'DATA_WIDTH | ID_Rs1_Data |
| Output | 'DATA_WIDTH | ID_Rs2_Data |

3.4 Immediate Generator (ImmGen)

The ImmGen module is responsible for extracting the immediate value from instructions and performing sign-extension based on the instruction type, generating a 32-bit immediate value.

Table 3.4: ImmGen I/O Ports

| I/O | Width | Port |
|--------|--------------|----------|
| Input | 'INSTR_WIDTH | ID_Instr |
| Input | 3 | Imm_Type |
| Output | 'DATA_WIDTH | ID_Imm |

3.5 Processor Controller

The Processor Controller is the CPU's control unit, responsible for decoding instructions and generating all control signals to coordinate the operation of each module.

Table 3.5: Processor Controller I/O Ports

| I/O | Width | Port |
|--------|---------------|------------------|
| Input | 'OPCODE_WIDTH | Opcode |
| Input | 1 | Branch_Taken |
| Input | 1 | ID_EX_Jump |
| Input | 1 | EX_Predict_Taken |
| Input | 'PC_WIDTH | ID_PC |
| Input | 'PC_WIDTH | Branch_PC |
| Output | 3 | Imm_Type |
| Output | 2 | ALU_op |
| Output | 2 | WB_sel |
| Output | 1 | Reg_w |
| Output | 1 | ALU_src1 |
| Output | 1 | ALU_src2 |
| Output | 1 | Mem_w |
| Output | 1 | Mem_r |
| Output | 1 | Branch |
| Output | 1 | Jump |
| Output | 1 | CSR_en |
| Output | 1 | IF_ID_Flush |
| Output | 1 | ID_EX_Flush_1 |

3.6 Arithmetic Logic Unit (ALU)

The ALU performs all arithmetic and logic operations, serving as the CPU's core computational unit.

Table 3.6: ALU I/O Ports

| I/O | Width | Port |
|--------|-------------|-------------|
| Input | 'DATA_WIDTH | Src1 |
| Input | 'DATA_WIDTH | Src2 |
| Input | 5 | ALU_CTRL_op |
| Output | 'DATA_WIDTH | ALU_Result |
| Output | 1 | Zero_Flag |

3.7 ALU Controller

The ALU Controller is responsible for combining the instruction's funct3 and funct7 fields with the Processor Controller's ALU_op signal to generate the actual operation control signals for the ALU.

Table 3.7: ALU Controller I/O Ports

| I/O | Width | Port |
|--------|-------|-------------|
| Input | 2 | ALU_op |
| Input | 3 | Funct3 |
| Input | 7 | Funct7 |
| Output | 4 | Mem_W_Strb |
| Output | 5 | ALU_CTRL_op |

3.8 Branch Processing Unit (BPU)

The BPU is responsible for handling the decision logic for all branch and jump instructions, determining whether to take a branch or jump.

Table 3.8: BPU I/O Ports

| I/O | Width | Port |
|--------|-------------|---------------|
| Input | 1 | ALU_Result[0] |
| Input | 1 | Zero_Flag |
| Input | 7 | Funct3 |
| Input | 1 | EX_Branch |
| Input | 'PC_WIDTH | EX_PC |
| Input | 'DATA_WIDTH | EX_Imm |
| Output | 'PC_WIDTH | PC_Plus_Imm |
| Output | 1 | Branch_Taken |

3.9 Load Data Unit (LDU)

The LDU is responsible for processing data read from memory by Load instructions, performing appropriate data extraction and sign-extension based on different Load instruction types.

Table 3.9: LDU I/O Ports

| I/O | Width | Port |
|--------|-------------|------------|
| Input | 3 | Funct3 |
| Input | 'DATA_WIDTH | Mem_R_Data |
| Output | 'DATA_WIDTH | LDU_Result |

3.10 Data Memory

The Data Memory module provides byte-addressable read and write access for the MEM stage. It supports byte, half-word, and word granularity through write strobe control, and adopts a write-first policy for simultaneous read/write access to the same address.

Table 3.10: Data Memory I/O Ports

| I/O | Width | Port |
|--------|----------------------|------------|
| Input | 1 | CLK |
| Input | 1 | Mem_r |
| Input | 1 | Mem_w |
| Input | 'DATA_MEM_ADDR_WIDTH | Mem_Addr |
| Input | 'DATA_MEM_WIDTH | Mem_W_Data |
| Input | 4 | Mem_W_Strb |
| Output | 'DATA_MEM_WIDTH | Mem_R_Data |

3.11 Five-Stage Pipelined System

This project utilizes a Five-Stage Pipelined System to partition the entire Processor into five stages: IF, ID, EX, MEM, and WB, improving instruction execution efficiency. If you want to see more simple implementation about Five-Stage Pipelined System, you can go to my [Github Repo of 2025 NTUST Computer Organization Course](#). The Repo implement Five-Stage Pipelined MIPS CPU step by step.

3.12 Forwarding, Hazard and Flush Detection Unit

This module is responsible for detecting and resolving various hazard conditions in the pipeline, ensuring correct pipeline execution and maintaining data consistency.

3.12.1 Forwarding Detection

Detects **Data Hazards** in the EX and MEM stages, generating forwarding signals to deliver correct data to the ALU in advance, avoiding unnecessary stalls.

Table 3.11: Forwarding Detection I/O Ports

| I/O | Width | Port |
|--------|-------------|-------------|
| Input | 'ADDR_WIDTH | MEM_Rd_Addr |
| Input | 1 | MEM_Reg_w |
| Input | 'ADDR_WIDTH | WB_Rd_Addr |
| Input | 1 | WB_Reg_w |
| Input | 'ADDR_WIDTH | EX_Rs1_Addr |
| Input | 'ADDR_WIDTH | EX_Rs2_Addr |
| Output | 2 | Forward_A |
| Output | 2 | Forward_B |

```
1 wire load_EX_MEM = (MEM_Reg_w&&MEM_Rd_Addr!=0);
2 wire load_MEMORY_WB = (WB_Reg_w&&WB_Rd_Addr!=0);
3
4 always @(*) begin
5     if(load_EX_MEM && (MEM_Rd_Addr==EX_Rs1_Addr)) Forward_A = 2'b10; // ALU_Result
6     else if(load_MEMORY_WB && (WB_Rd_Addr==EX_Rs1_Addr)) Forward_A = 2'b01; // WB_DATA
7     else Forward_A = 2'b00;
8
9     if(load_EX_MEM && (MEM_Rd_Addr==EX_Rs2_Addr)) Forward_B = 2'b10; // ALU_Result
10    else if(load_MEMORY_WB && (WB_Rd_Addr==EX_Rs2_Addr)) Forward_B = 2'b01; // WB_DATA
11    else Forward_B = 2'b00;
12 end
```

Listing 3.1: Forwarding Detection Logic

3.12.2 Hazard Detection

Detects **Load-Use Hazards**. When a Load instruction is immediately followed by an instruction that uses the loaded data, it generates a stall signal to pause the pipeline for one cycle.

Table 3.12: Hazard Detection I/O Ports

| I/O | Width | Port |
|--------|-------------|---------------|
| Input | 'ADDR_WIDTH | Rs1_Addr |
| Input | 'ADDR_WIDTH | Rs2_Addr |
| Input | 'ADDR_WIDTH | Rd_Addr |
| Input | 1 | EX_Mem_r |
| Output | 1 | IF_ID_w |
| Output | 1 | ID_EX_Flush_0 |

```

1 always @(*) begin
2   if(EX_Mem_r&&((RdAddr==Rs1Addr) || (RdAddr==Rs2Addr))) begin
3     IF_ID_w = 1'b0;
4     ID_EX_Flush_0 = 1'b1;
5   end
6   else begin
7     IF_ID_w = 1'b1;
8     ID_EX_Flush_0 = 1'b0;
9   end
10 end

```

Listing 3.2: Hazard Detection Logic

3.12.3 Flush Detection

Detects branch/jump and control flow changes (**Control Hazards**), generating flush signals to convert invalid instructions in the IF and ID stages into bubbles (NOPs).

```

1 always @(*) begin
2   if(ID_PC == Branch_PC && (Branch_Taken || ID_EX_Jump)) IF_ID_Flush = 0;
3   else begin
4     if(Branch_Taken || ID_EX_Jump) IF_ID_Flush = 1;
5     else if(EX_Predict_Taken) IF_ID_Flush = 1;
6     else IF_ID_Flush = 0;
7   end
8 end
9
10 assign ID_EX_Flush_1 = (ID_PC == Branch_PC && (Branch_Taken || ID_EX_Jump)) ? 0 :
11   ((Branch_Taken || ID_EX_Jump) || (EX_Predict_Taken && ~(
      Branch_Taken || ID_EX_Jump))) ;

```

Listing 3.3: Flush Detection Logic

3.13 Branch History Table (BHT) & Branch Tag Buffer (BTB)

The BHT and BTB together implement a dynamic branch prediction mechanism. By recording historical branch behavior to predict branch outcomes, they reduce pipeline stalls caused by control hazards.

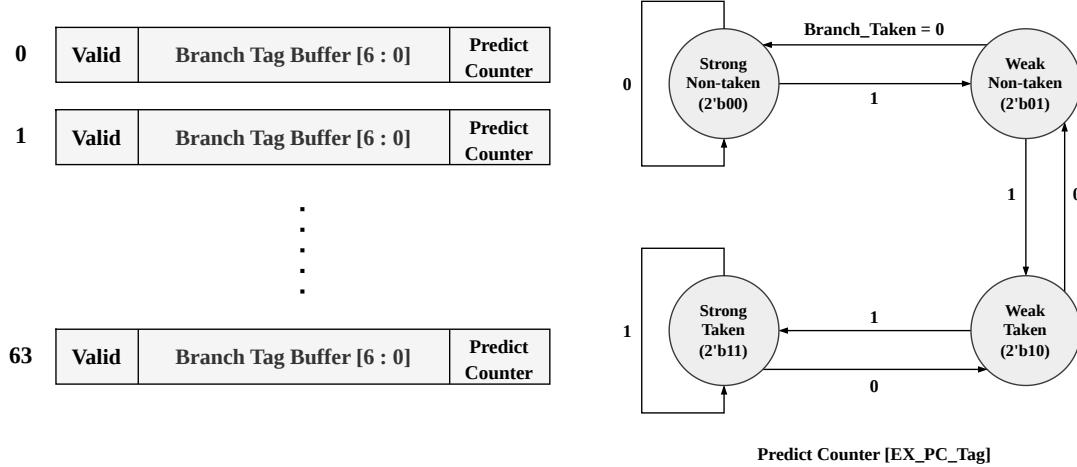


Figure 3.2: BHT & BTB Dynamic Branch Prediction Architecture

3.13.1 BHT (Branch History Table)

Uses 2-bit saturating predict counters to record the historical behavior of branch instructions and predict whether branches will be taken.

- Employs four states: Strong Non-Taken, Weak Non-Taken, Weak Taken, Strong Taken
- The prediction counter is updated at the **EX stage**, once the actual branch outcome (Branch_Taken) is resolved, ensuring the state machine reflects the most recent branch behaviour
- Uses lower bits of the PC as an index (PC Tag) to access the prediction table

Table 3.13: BHT I/O Ports

| I/O | Width | Port |
|--------|---------------|--------------|
| Input | 1 | CLK |
| Input | 1 | RST_N |
| Input | 'BHT_PC_WIDTH | PC_Tag |
| Input | 1 | Branch_Taken |
| Input | 'BHT_PC_WIDTH | EX_PC_Tag |
| Output | 1 | Predict |

3.13.2 BTB (Branch Target Buffer)

Stores the target addresses of branch instructions, providing the jump target when a branch is predicted taken.

- Records the PC tag of branch instructions and their corresponding target addresses
- Supports fast lookup to avoid waiting for branch target calculation
- Works in conjunction with BHT prediction results to update the PC in the IF stage

Table 3.14: BTB I/O Ports

| I/O | Width | Port |
|--------|---------------|--------------|
| Input | 1 | CLK |
| Input | 1 | RST_N |
| Input | 'BHT_PC_WIDTH | PC_Tag |
| Input | 1 | Branch_Taken |
| Input | 'BHT_PC_WIDTH | EX_PC_Tag |
| Input | 'PC_WIDTH | Branch_PC |
| Output | 'PC_WIDTH | BTB_PC |
| Output | 1 | BTB_Valid |

3.14 Control State Register (CSR)

The CSR module implements RISC-V Control and Status Registers, providing system-level state management, exception handling, and privilege mode control functionality.

Table 3.15: CSR I/O Ports

| I/O | Width | Port |
|--------|-------------|------------|
| Input | 1 | CLK |
| Input | 1 | RST_N |
| Input | 1 | CSR_en |
| Input | 12 | CSR_Addr |
| Input | 'DATA_WIDTH | CSR_W_Data |
| Input | 3 | Funct3 |
| Output | 'DATA_WIDTH | CSR_R_Data |

3.14.1 Supported CSR Registers

Table 3.16: Supported CSR Registers

| Name | Address | Description | Access |
|---------|---------|-----------------------------------|-----------------------|
| mstatus | 0x300 | Machine Status Register | CSRRW / CSRRS / CSRRC |
| mtvec | 0x305 | Machine Trap Vector Base Address | CSRRW / CSRRS / CSRRC |
| mepc | 0x341 | Machine Exception Program Counter | CSRRW / CSRRS / CSRRC |
| mcause | 0x342 | Machine Exception Cause | CSRRW / CSRRS / CSRRC |
| rdcycle | 0xC00 | Cycle Counter | Read-only |

3.15 Verification

The verification approach for this project does not use the [Spike RISC-V ISA Simulator](#). Instead, a faster solution is adopted. Using Claude Code, a RISC-V CPU behavioral model [Golden_Result.py](#) is rapidly built in Python to simulate instruction execution and generate expected outputs for the Register File and Data Memory as the Golden Result.

The Golden Result is then automatically compared with RF.out and DM.out generated by the Verilog Testbench to verify the correctness of the RTL design.

This project includes [12 TestCases](#) designed to verify that all RV32I and RV32M instructions are correctly implemented. These test cases cover different types of operations, including arithmetic and logic operations, memory access, branches and jumps, as well as multiply-divide extension instructions, ensuring the CPU executes correctly under various instruction combinations and boundary conditions.

Each TestCase is written in assembly language and then converted to machine code via [Instr_Transfer.py](#) to be written into IM.dat. Users can complete the workflow of selecting a Test Case → conversion → result verification through the automated [Verify_Script.py](#) script.

Chapter 4

AXI-4 Lite Bus, BRAM and Cache

4.1 AXI-4 Lite Interface with BRAM

AXI (Advanced eXtensible Interface) **[AXI Protocol Specification]** is a high-performance, high-bandwidth, and low-latency on-chip interconnect protocol within the AMBA standard. Its architectural core lies in the decoupling of address/control and data phases, enabling maximum transmission efficiency and optimized bus utilization. In this implementation, an AXI4-Lite bus will be implemented to interface with the BRAM for data storage and retrieval.

4.1.1 Required signals on an AXI4-Lite interface

Table 4.1: Required signals on an AXI4-Lite interface

| Name | Width | Description |
|---------|-------|--|
| AWVALID | 1 | Asserted high to indicate that the signals on the AW channel are valid. |
| AWREADY | 1 | Asserted high to indicate that a transfer on the AW channel can be accepted. |
| WVALID | 1 | Asserted high to indicate that the signals on the W channel are valid. |
| WREADY | 1 | Asserted high to indicate that a transfer on the W channel can be accepted. |
| WSTRB | 4 | Byte strobe signals indicating which byte lanes of WDATA contain valid data. |
| BVALID | 1 | Asserted high to indicate that the signals on the B channel are valid. |
| BREADY | 1 | Asserted high to indicate that a transfer on the B channel can be accepted. |
| BRESP | 2 | Write response status indicating the result of the write transaction. |
| ARVALID | 1 | Asserted high to indicate that the signals on the AR channel are valid. |
| ARREADY | 1 | Asserted high to indicate that a transfer on the AR channel can be accepted. |
| RVALID | 1 | Asserted high to indicate that the signals on the R channel are valid. |
| RREADY | 1 | Asserted high to indicate that a transfer on the R channel can be accepted. |
| RRESP | 2 | Read response status indicating the result of the read transaction. |

Note: BRESP and RRESP are fixed to 2'b00 (OKAY) in this implementation.

4.1.2 AXI4-Lite interface to BRAM

In this implementation, For Single-Port RAM with simultaneous read and write to the same address, a **write-first mode** is adopted and reading from BRAM requires 1 cycle.

`Tested.v` is the top-level module that connects `AXI4_Lite_Bus.v` with the Single-Port BRAM IP generated by Vivado. `Pattern.v` serves as the testbench, testing four different cases:

- **Test Case 1 :** Sequential Write/Read with AW/W separate (AW and W channels arrive in separate cycles)
- **Test Case 2 :** Parallel Write/Read with AW/W together (AW and W channels arrive in same cycle)
- **Test Case 3 :** Byte-Enable Write (W arrives before AW, W_STRB = 4'b0011)
- **Test Case 4 :** Sequential Pattern Write then Read (Write sequential pattern, then read back all data)

4.2 Two-way set-associative Cache with AXI4 Lite to BRAM

4.2.1 Instruction Cache (Only Read Cache)

In this implementation, the Instruction Memory is designated as **read-only**. Consequently, the Instruction Cache is also read-only and does not support data write-backs or write-throughs. The AXI4-Lite bus interface remains unchanged, as write functionality is disabled by simply masking all write-related signals.

When the CPU requires an instruction, it issues CPU_REQ along with CPU_REQ_ADDR. Once the instruction is ready, the system asserts CPU_REQ_VALID, at which point the CPU retrieves the instruction from the CPU_REQ_DATA bus. The full RTL design and verification for this module is available [here](#).

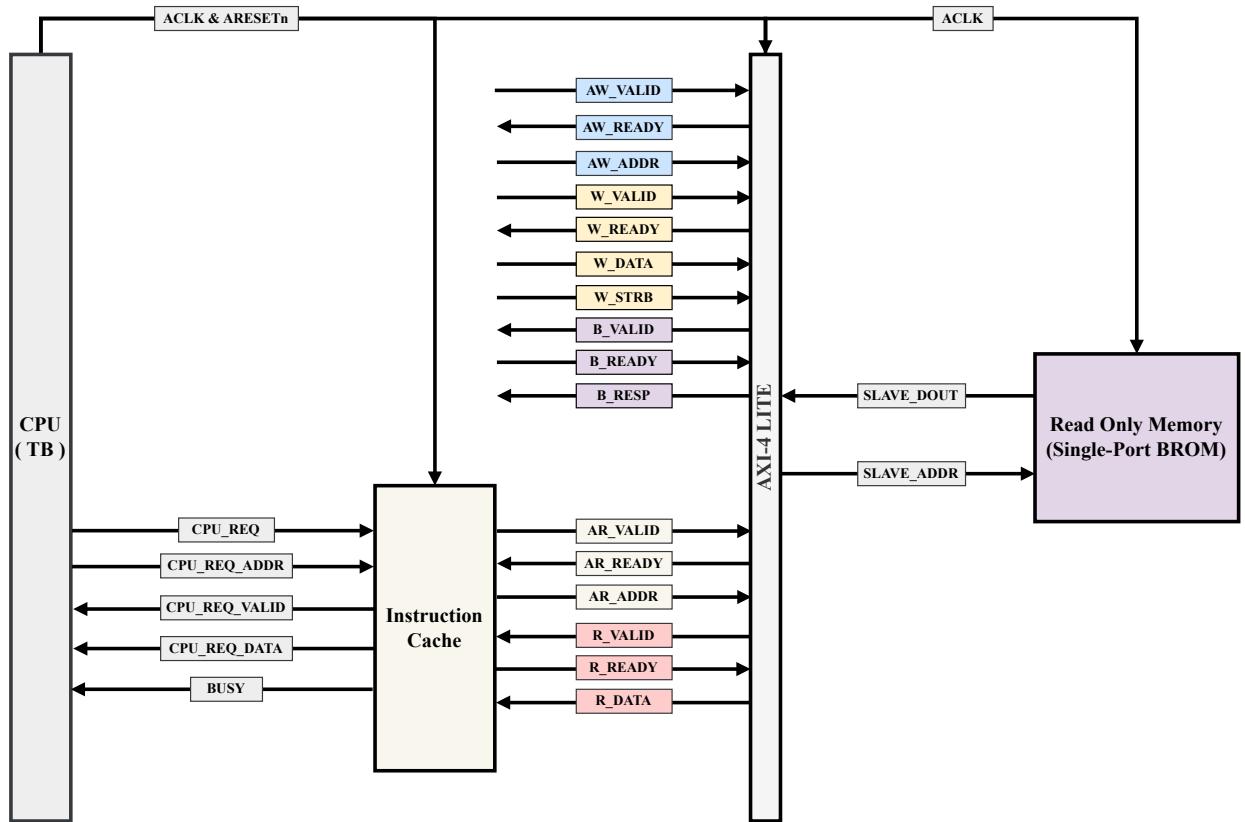


Figure 4.1: Two-way set-associative Cache with AXI4-Lite to BRAM

4.2.2 Data Cache

In this implementation, the Data Cache employs a Write-Through policy instead of Write-Back with a Write Buffer, which presents an opportunity for future optimization. The Data Cache read operation follows the same logic as the Instruction Cache described in [Section 4.2.1](#). It is important to note that the write and read operations share the same ADDR PORT, and the system follows a Write-First approach similar to Memory. The full RTL design and verification for this module is available [here](#).

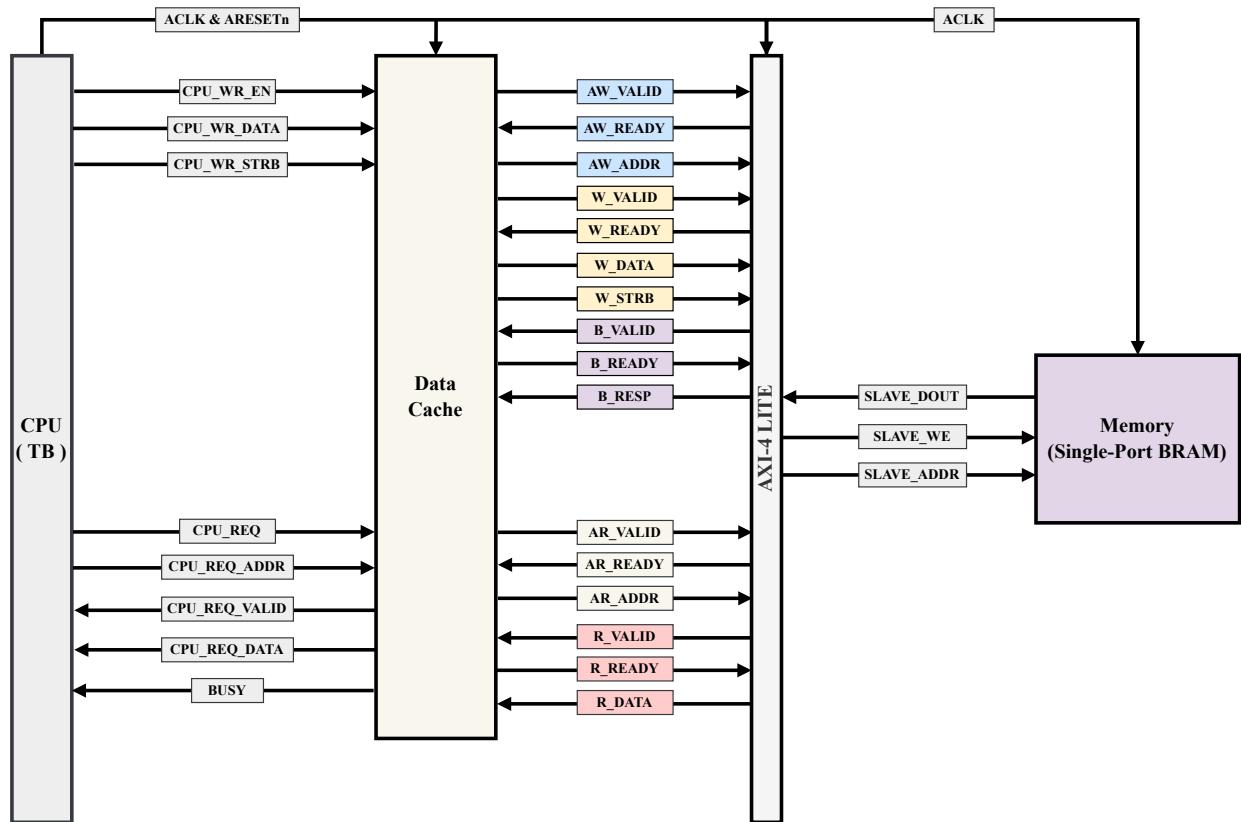


Figure 4.2: Data Cache with AXI4-Lite to BRAM

Chapter 5

RISC-V Processor

5.1 Full System Integration

In previous chapters, we successfully implemented a 5-stage pipelined RISC-V CPU supporting the RV32I and RV32M instruction sets, an AXI4-Lite Bus Interface with BRAM, an I-Cache (AXI4-Lite to BROM), and a D-Cache (AXI4-Lite to BRAM). In this chapter, we will integrate both caches into the CPU to complete the **Final RISC-V Processor**, as you can see the Figure 5.1.

It is worth noting that in this chapter, some of the modules introduced in Chapter 3 will be updated with additional control signals or logic to accommodate the cache integration. However, the overall architecture remains largely unchanged.

The integration of the I-Cache is relatively straightforward : while waiting for a valid instruction read (a cache miss or pending fetch), we simply **disable write-enable for the IF/ID pipeline registers** and **stall the Program Counter (PC)** to prevent it from updating. This effectively freezes the front-end until the instruction is ready.

Integrating the D-Cache is considerably more complex. Because memory read and write operations spend multiple clock cycles, the entire processor must stall until the transaction completes successfully. To achieve this, **stall signals must be asserted across all pipeline stages**.

Furthermore, extra care must be taken to ensure that a **Branch_Taken signal is not triggered during a stall**, which prevents illegal state transitions or incorrect instruction execution. We can determine the completion of a write operation by monitoring the AXI-4 Lite response via **BVALID && BREADY**. For read operations, we track the status through **CPU_REQ_VALID**. Ultimately, whenever the D-Cache is in a **busy state**, the entire pipeline must remain stalled.

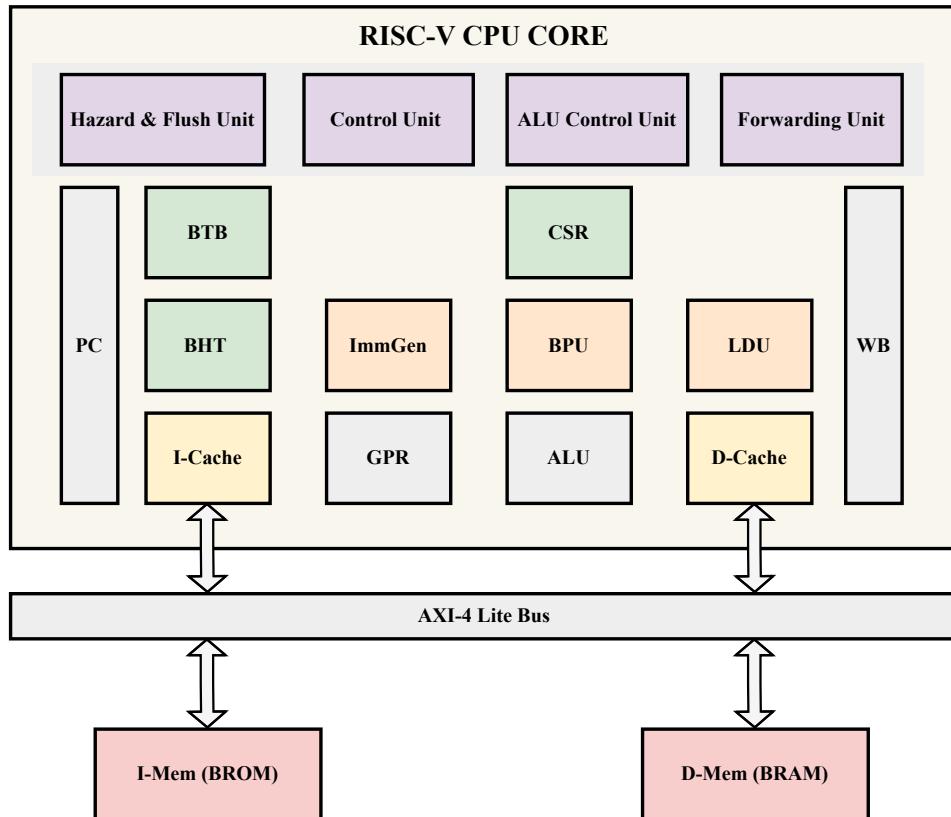


Figure 5.1: RISC-V Processor Overview

5.2 System-Level Verification

The verification process remains consistent with the flow established in Section 3.15, with a few key additions to the automation suite. We have updated `Verify_Script.py` to include a script that converts IM.dat into a .coe format, specifically for the Instruction Memory (BROM IP).

Furthermore, to facilitate the rapid export of Data Memory content via the TestBench, we opted not to use the Vivado-generated BRAM IP. Instead, we implemented a custom module, `D_BRAM.v`, which replicates the IP's functionality while allowing for easier file I/O operations and data extraction.

More importantly, to accelerate the verification process, I have written the Vivado simulation TCL commands into a `Script.tcl` file and integrated it with the `Verify_Script.py` script to enable automated RTL simulation and verification. Users can choose to run a single Test Case or execute all Test Cases at once.

```
ALL TESTS PASSED!
Your CPU simulation matches the golden reference perfectly.

=====
Final Summary - All Test Cases
=====



| TestCase   | RF   | DM   | Result |
|------------|------|------|--------|
| TestCase1  | PASS | PASS | ✓ PASS |
| TestCase2  | PASS | PASS | ✓ PASS |
| TestCase3  | PASS | PASS | ✓ PASS |
| TestCase4  | PASS | PASS | ✓ PASS |
| TestCase5  | PASS | PASS | ✓ PASS |
| TestCase6  | PASS | PASS | ✓ PASS |
| TestCase7  | PASS | PASS | ✓ PASS |
| TestCase8  | PASS | PASS | ✓ PASS |
| TestCase9  | PASS | PASS | ✓ PASS |
| TestCase10 | PASS | PASS | ✓ PASS |
| TestCase11 | PASS | PASS | ✓ PASS |
| TestCase12 | PASS | PASS | ✓ PASS |



=====

ALL 12 TEST CASES PASSED!
```

Figure 5.2: Verification Results

5.3 Synthesis Results

The design was synthesised targeting the Xilinx ZCU104 (UltraScale+) platform at a clock frequency of **250 MHz**. The table below summarises the key post-synthesis resource utilisation, followed by the complete cell-level breakdown reported by Vivado.

Table 5.1: Resource Utilisation Summary

| Resource | Count |
|------------------------|----------------|
| LUT (total) | 55,911 |
| Flip-Flop (total) | 76,491 |
| DSP48E2 | 12 |
| BRAM | 2 |
| Clock Frequency | 250 MHz |

Table 5.2: Vivado Report Cell Usage

| Cell | Count | Cell | Count |
|----------------------|--------|-------------------|--------|
| <i>Clock / Carry</i> | | <i>LUT</i> | |
| BUFG | 1 | LUT1 | 225 |
| CARRY8 | 717 | LUT2 | 664 |
| <i>DSP</i> | | LUT3 | 2,677 |
| DSP_ALU | 12 | LUT4 | 1,646 |
| DSP_A_B_DATA | 12 | LUT5 | 6,316 |
| DSP_C_DATA | 12 | LUT6 | 44,383 |
| DSP_MULTIPLIER | 12 | <i>MUX</i> | |
| DSP_M_DATA | 12 | MUXF7 | 18,036 |
| DSP_OUTPUT | 12 | MUXF8 | 8,343 |
| DSP_PREADD | 12 | <i>Memory</i> | |
| DSP_PREADD_DATA | 12 | RAMB18E2 | 1 |
| <i>Flip-Flop</i> | | RAMB36E2 | 1 |
| FDCE | 10,942 | <i>I/O Buffer</i> | |
| FDPE | 5 | IBUF | 2 |
| FDRE | 65,544 | OBUF | 32 |